

COMP 8567
Advanced Systems Programming

Pipes

Outline

- Unnamed Pipes
- The pipe() system call
- The dup2() system call
 - I/O redirection using dup2()
 - Reversal of I/o redirection using dup()
 - Implementing the shell piping mechanism using dup2()
- FIFOs or named pipes
- Summary

Unnamed Pipes

- Unnamed Pipes, known as pipe, are a mechanism for inter-process communication.
- They are used by shells to connect one utility's **standard output** with the **standard input** of another utility.
 - **Example :** `$ ls|wc -w`
- Unnamed pipes are **in-memory files** created by the kernel. The **kernel provides the synchronization** between the processes accessing the same pipe
- Pipes are the oldest form of Unix IPC.
- Pipes have two limitations:
 - Data flows in one direction only
 - They can be **used only** between processes that have a **common ancestor** (related processes only)
 - Typically, a process creates a pipe, forks, and then uses the pipe to exchange information with its child.

Unnamed Pipes..


- Created In-memory (main memory) by the kernel
- Lasts as long as the process that created it lives.
- **Does not have a name**

The pipe() system call

Synopsis : **int pipe(int fd[2])**

returns 0 when successful and -1 otherwise

pipe() creates a pipe and returns **two file descriptors** fd[0] and fd[1], where fd[0] is open for reading and fd[1] is open for writing.



→ A pipe is a one-way communication channel between two **related processes**

ex1.c //The same program writes into and reads from a pipe


```
// C program to illustrate pipe() system call
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
char* msg1 = "Welcome to COMP 8567\n";
//As a simple example, the same process writes and reads from the pipe
int main()
{
    char inbuf[20];
    int p[2], i;
    if (pipe(p) < 0) //Invoke the pipe() system call
        exit(1);
    //write msg1 into the pipe using the write FD p[1]
    write(p[1], msg1, 20);
    //read the contents of the pipe into inbuf using the read FD p[0]
    int n= read(p[0], inbuf, 20);
    printf("The cotents of the pipe are\n%s", inbuf);
    printf("\nThe number of bytes read were %d\n",n);
    return 0;
}
```

When reading from or writing to a pipe, the following rules apply :

1. If a process reads from an empty pipe whose write fd is still open, it sleeps until some input becomes available.
 2. If a process tries to read from a pipe more bytes than are present, read() reads **all available bytes** and returns the number of bytes read.
 3. If a process writes to a pipe whose **read fd has been closed**, the write operation fails and the writer process receives a **SIGPIPE** (//Default action: Terminate)
- Note : In case of multiple processes writing to the same pipe, a write of up to PIPE_BUF bytes is guaranteed to be atomic
 - Data from different writer processes will not be interleaved.

pipe() followed by fork()

All the descendent processes get a copy of the file descriptors and can use them to access the pipe (to both read and write from and into the pipe)




```

#include <stdio.h> //ex2d.c
//Pipe between parent and child process
int main(int argc, char *argv[]){
int fd[2];
int k=pipe(fd);//create a pipe
if(k == -1)
exit(1);

int pid=fork();

if(pid>0)// Parent Process
{
close(fd[0]);//since parent does not use fd[0]
char *message="Hello child process!";
int n=write(fd[1], message, 40);
printf("\nThe number of characters written were %d\n",n);
}

```


```

else
{
char ch;
//char buff1[255];
char *buff1;
close(fd[1]);//since child does not use fd[1]
printf("Child Process: Parent has sent the following message:\n");
int n=read(fd[0],buff1,40);
printf("\n%s",buff1);
printf("\nThe number of characters read were %d\n",n);

}
exit(0);
}

```

I/O Redirection using dup2() system call



```

//redirect.c
// Demonstrates I/O redirection with dup2()

int main(void) {

    int number1, number2, sum;
    int fdx = open("filex.txt", O_RDONLY); //scanf reads from this file
    int fdy = open("filey.txt", O_RDWR); //printf writes into this file

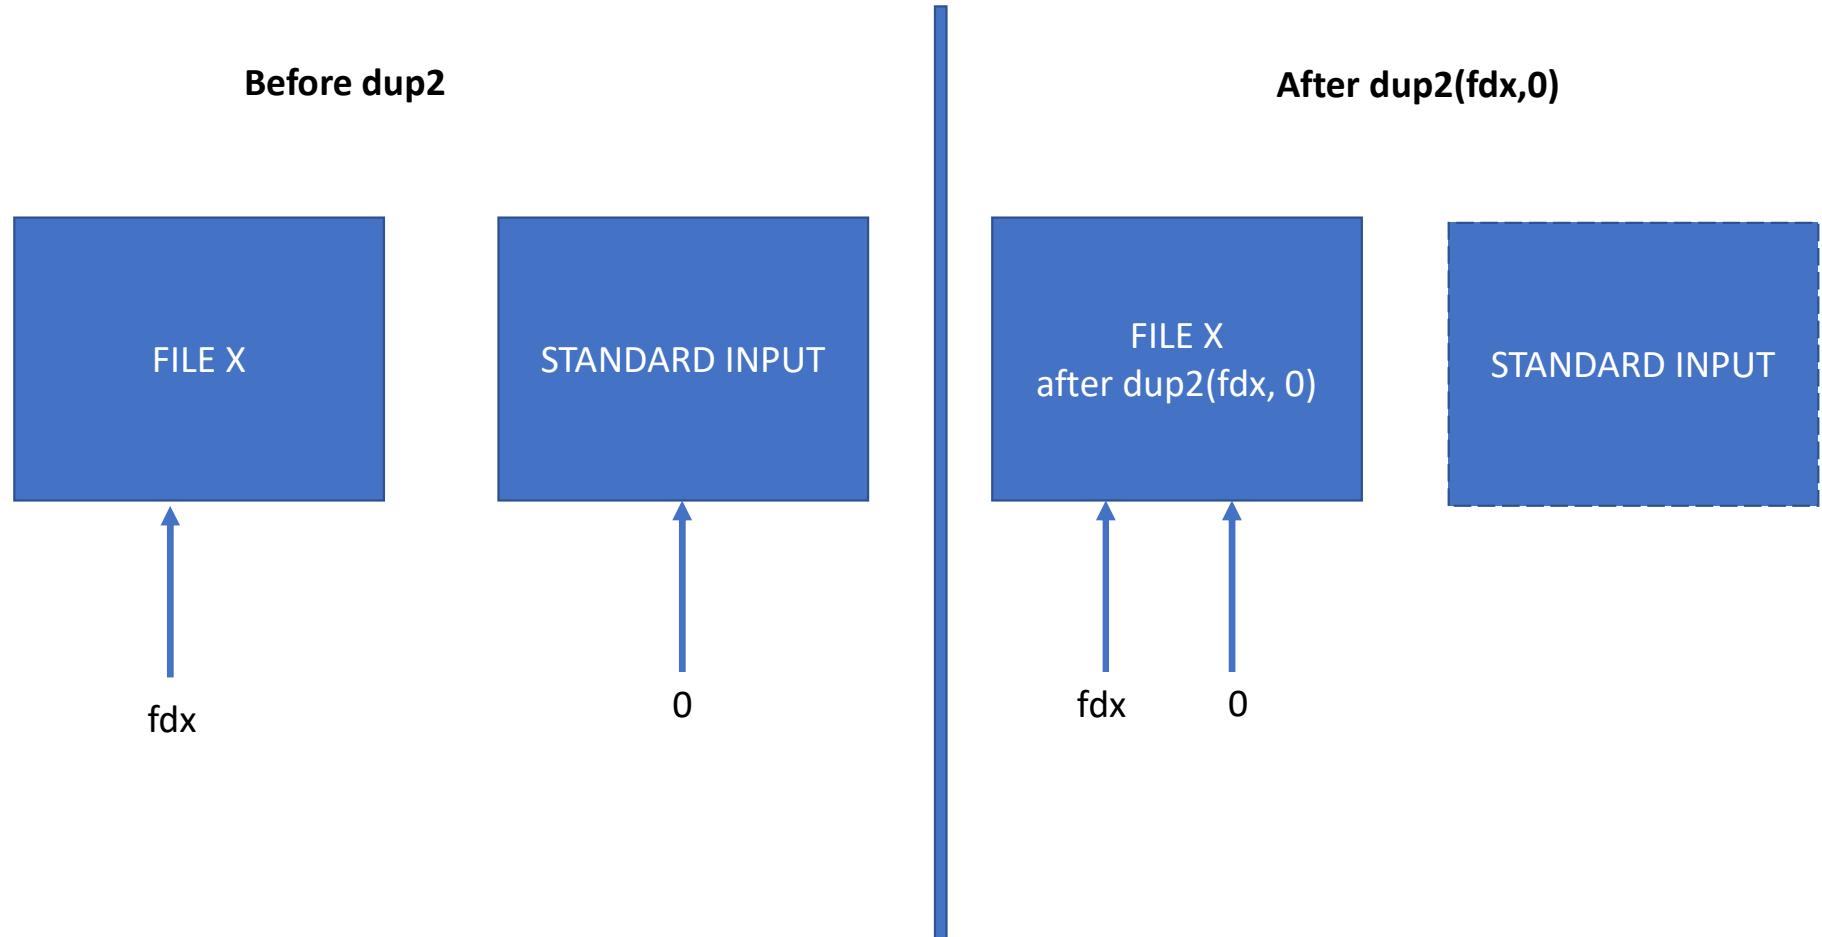
    int ret1= dup2(fdx,0); //standard input
    if(ret1 < 0) {
        printf("Unable to duplicate the STDIN file descriptor.");
        exit(EXIT_FAILURE); }
    int ret2=dup2(fdy,1); // standard output
    if(ret2 < 0) {
        printf("Unable to duplicate the STDOUT file descriptor.");
        exit(EXIT_FAILURE);}

    scanf("%d %d", &number1, &number2); //reads number 1 and number2 from
    filex.txt and not the standard input
    sum = number1 + number2;
    printf("The sum of two numbers is\n"); // writes into filex.txt and not the std output
    printf("%d + %d = %d\n", number1, number2, sum); // writes into filex.txt
    return EXIT_SUCCESS;
}


```

I/O Redirection and Reversing using dup2 and dup System Calls
in Linux


Redirecting Standard Input



Redirecting Standard Output



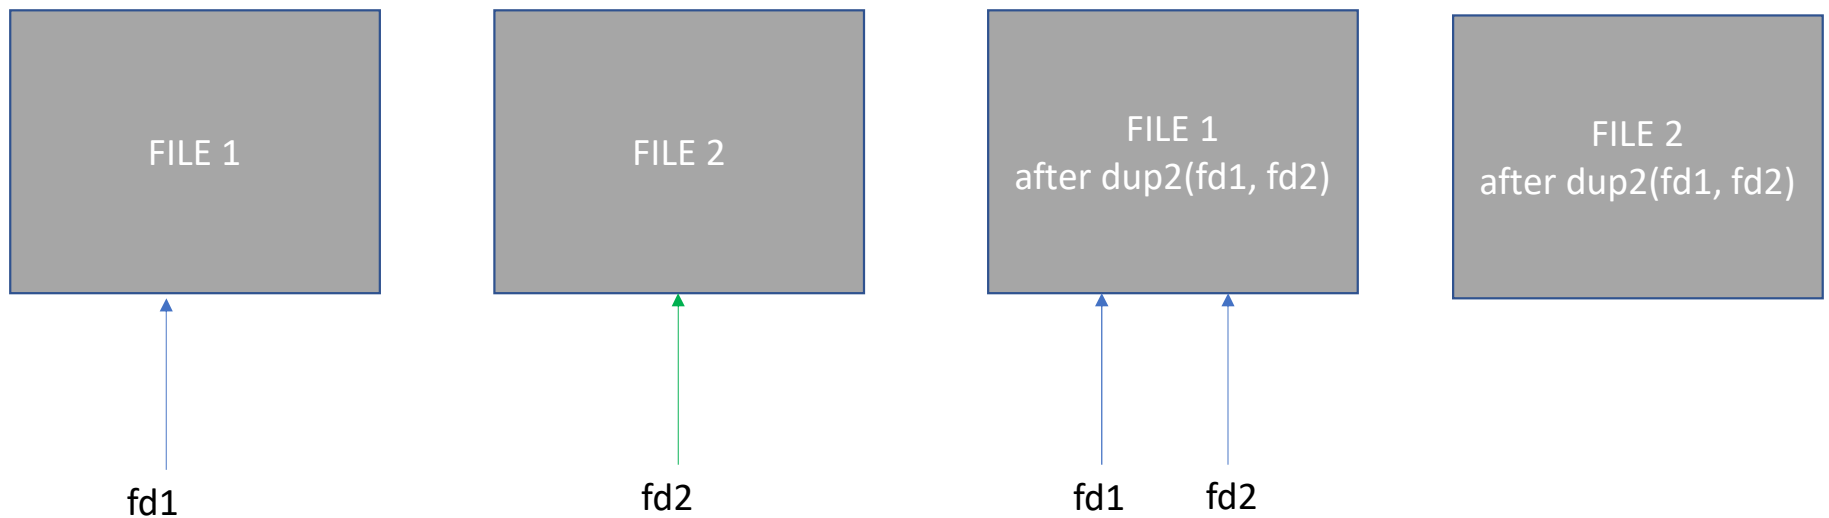
scanf and printf operations after redirection



The dup2() system call //

Synopsis : *int dup2(int fd1, int fd2);*

- The dup2() function causes the file descriptor fd2 to refer to the same file referred by fd1.
- The fd1 argument is a file descriptor referring to an **open file**
- fd2 is a non-negative integer less than FOPEN_MAX (0 to FOPEN_MAX)
- On success, returns the value of *fd2*
- On failure, returns -1



dup2() effect

int **dup2**(*fd1*,*fd2*)

- On success, returns the value of *fd2*
- On failure, returns -1


Effect: After a successful dup2 call, *fd2*

will refer to the file referred by *fd1*

// *fd1* should be a valid open file descriptor

int dup2(fd1,fd2) Rules

RULE 1: If fd2 already refers to an open file that is not fd1, fd2 is closed first (and then assigned to fd1)



RULE-2: If fd2 already refers to fd1, dup2(fd1,fd2) returns fd2

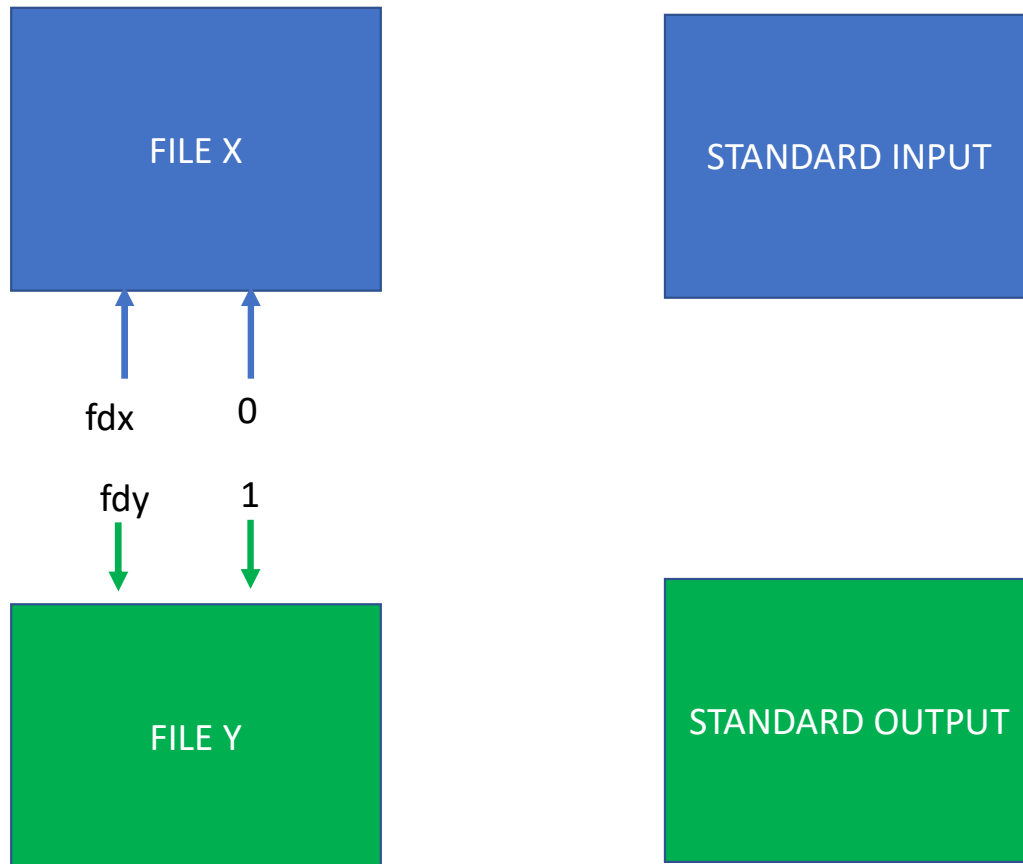


RULE-3: if fd1 is **not a valid open file** descriptor, dup2(fd1,fd2) returns -1 , **fd2 will not be closed**

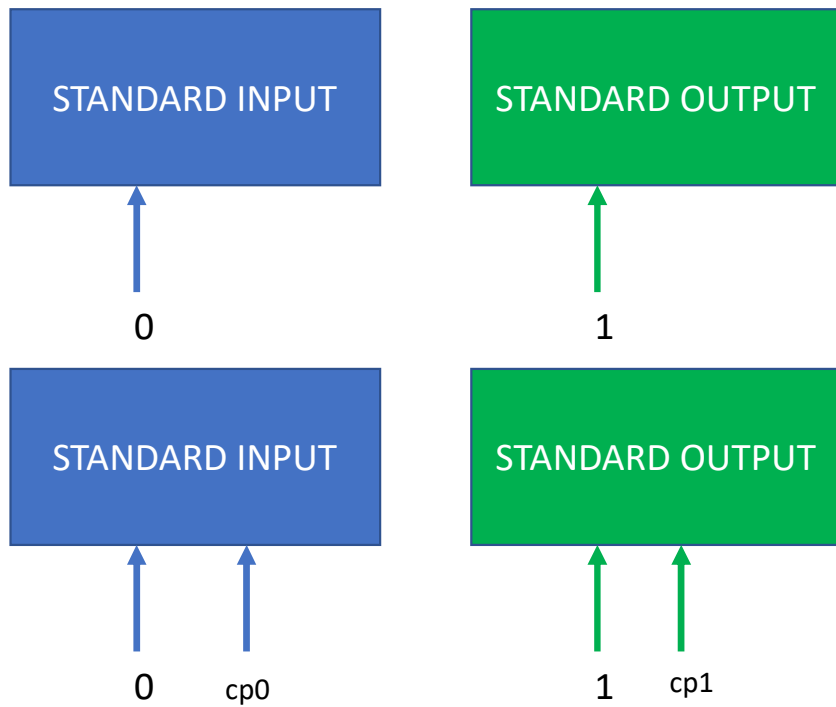


Question:

Can 0 and 1 go back to referring standard input and standard output again?




Reversal of Input/Output Redirection

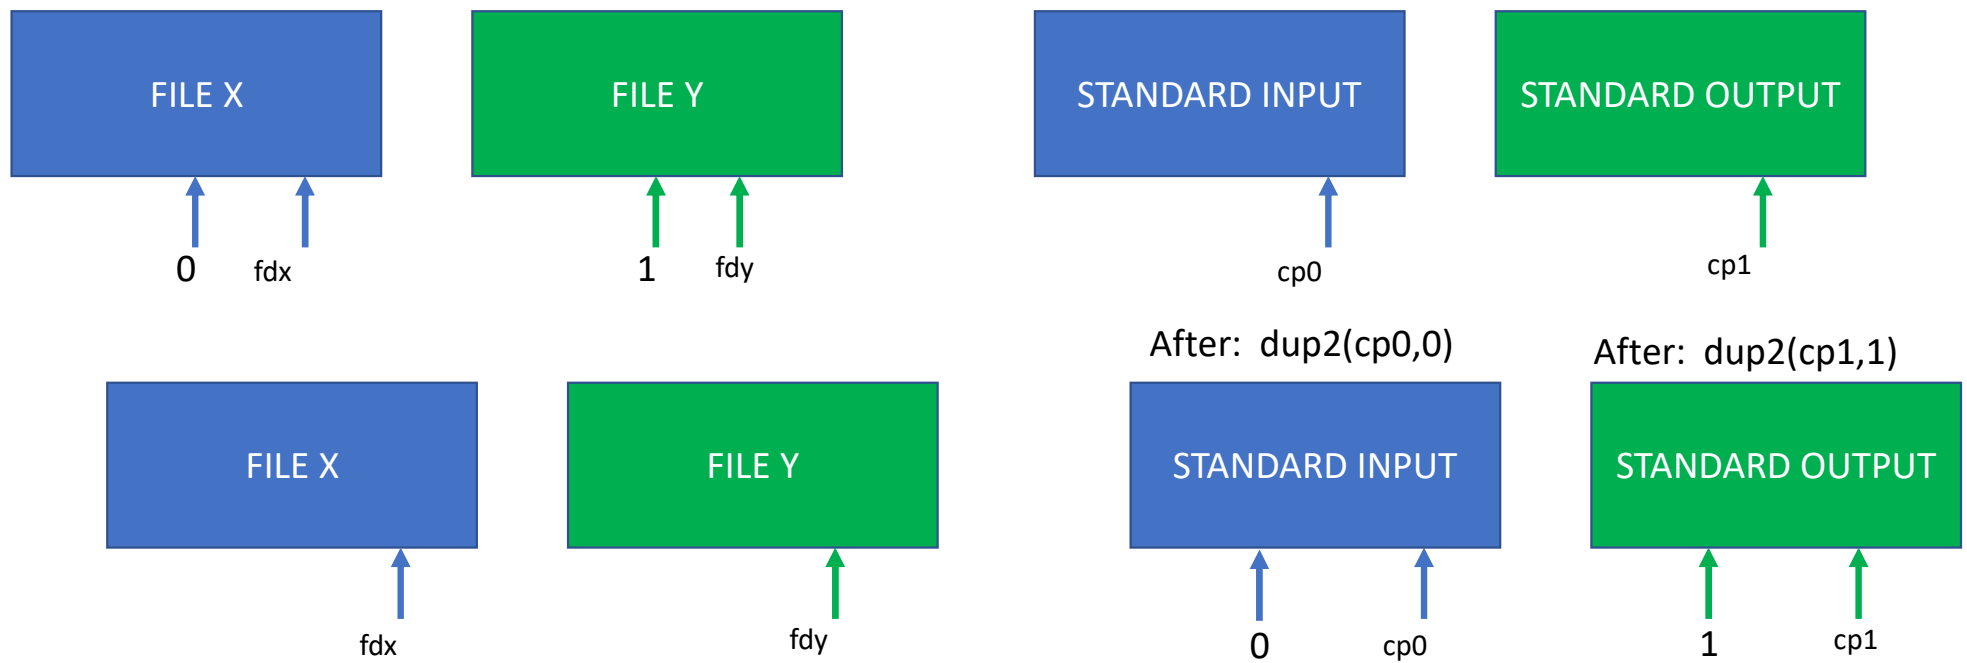
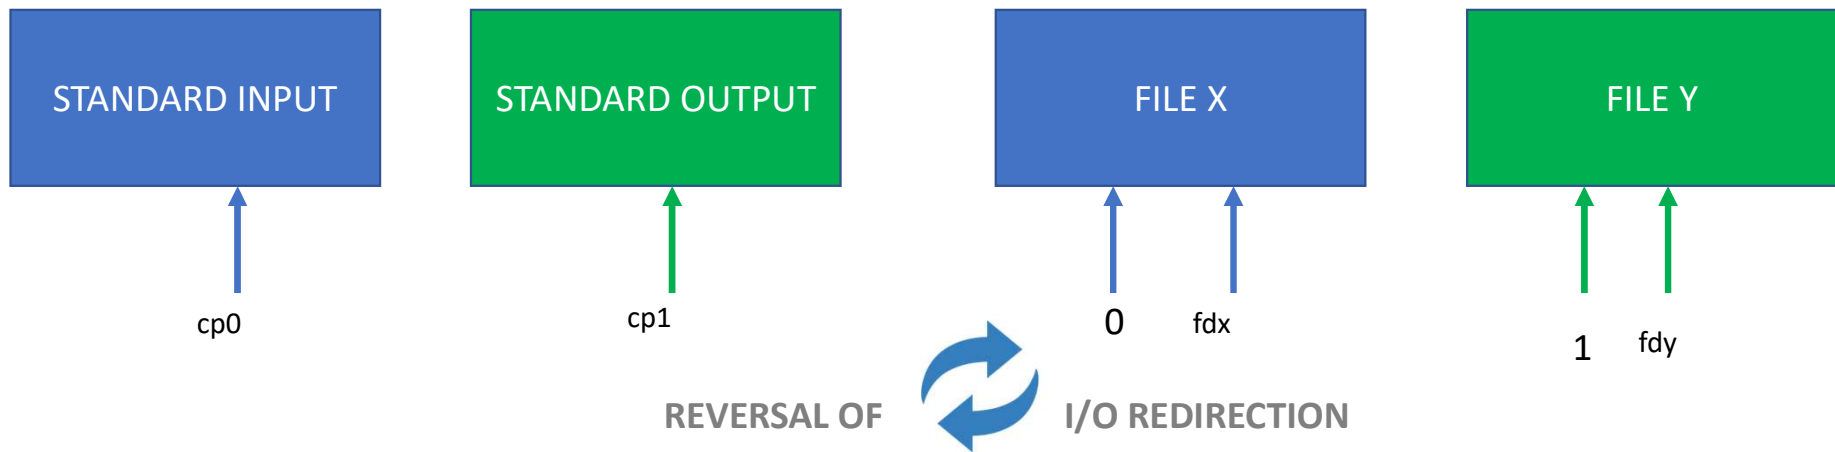


Creating copies of 0 and 1 before redirection using the dup() system call

After:
`cp0= dup(0);`
`cp1 = dup(1);`

I/O redirection after:
`dup2(fdx,0)`
`dup2(fdy,1)`





How to create a duplicate of STDIN and STDOUT before redirection?

Using **int dup(fd1)**

- `int fd2=dup(fd1);`
- `fd2= first unused` file descriptor and will now point to `fd1`

//guaranteed to be an unused fd

Ex:

```
int cp0= dup(0);
```

```
int cp1 = dup(1);
```

// Note: `dup2()` can also be used, but we will risk overwriting existing fds.

// `dup()` is therefore preferred.

```
dup2(0,11) //Not sure if 11 and 12 are already referring to open files
```

```
dup2(1,12)
```

```
#include <stdio.h> //backtoio.c
```

```
#include <stdlib.h>
```

```
#include <unistd.h>
```

```
#include <fcntl.h>
```

```
//Redirection and reversal
```

```
int main(void) {
```

```
int number1, number2, sum;
```

```
int fdx = open("filex.txt", O_RDONLY);
```

```
int fdy = open("filey.txt", O_RDWR);
```

```
int cp0=dup(0);
```

```
int cp1=dup(1);
```

```
int ret1= dup2(fdx,0); //0 is the fd of standard input
```

```
if(ret1 < 0) {
```

```
    printf("Unable to duplicate the STDIN file descriptor.");
```

```
    exit(EXIT_FAILURE);
```

```
}
```

```
int ret2= dup2(fdy, 1);
```

```
printf("\nThe value of ret2 is %d\n",ret2);
```

```
printf("\n Redirection of the standard output\n");
```

```
if(ret2 < 0) {
```

```
    printf("Unable to duplicate the STDOUT file descriptor.");
```

```
    exit(EXIT_FAILURE);
```

```
}
```

```
scanf("%d %d", &number1, &number2);
```

```
sum = number1 + number2;
```

```
printf("\nThe sum of two numbers is\n");
```

```
printf("%d + %d = %d\n", number1, number2, sum);
```

```
fflush(stdout);
```

```
// Reversal of Redirection
```

```
int retval1= dup2(cp0,0);
```

```
int retval2=dup2(cp1,1);
```

```
printf("\nBack to standard input, enter the value of num\n");
```

```
fflush(stdout);
```

```
fflush(stdin);
```

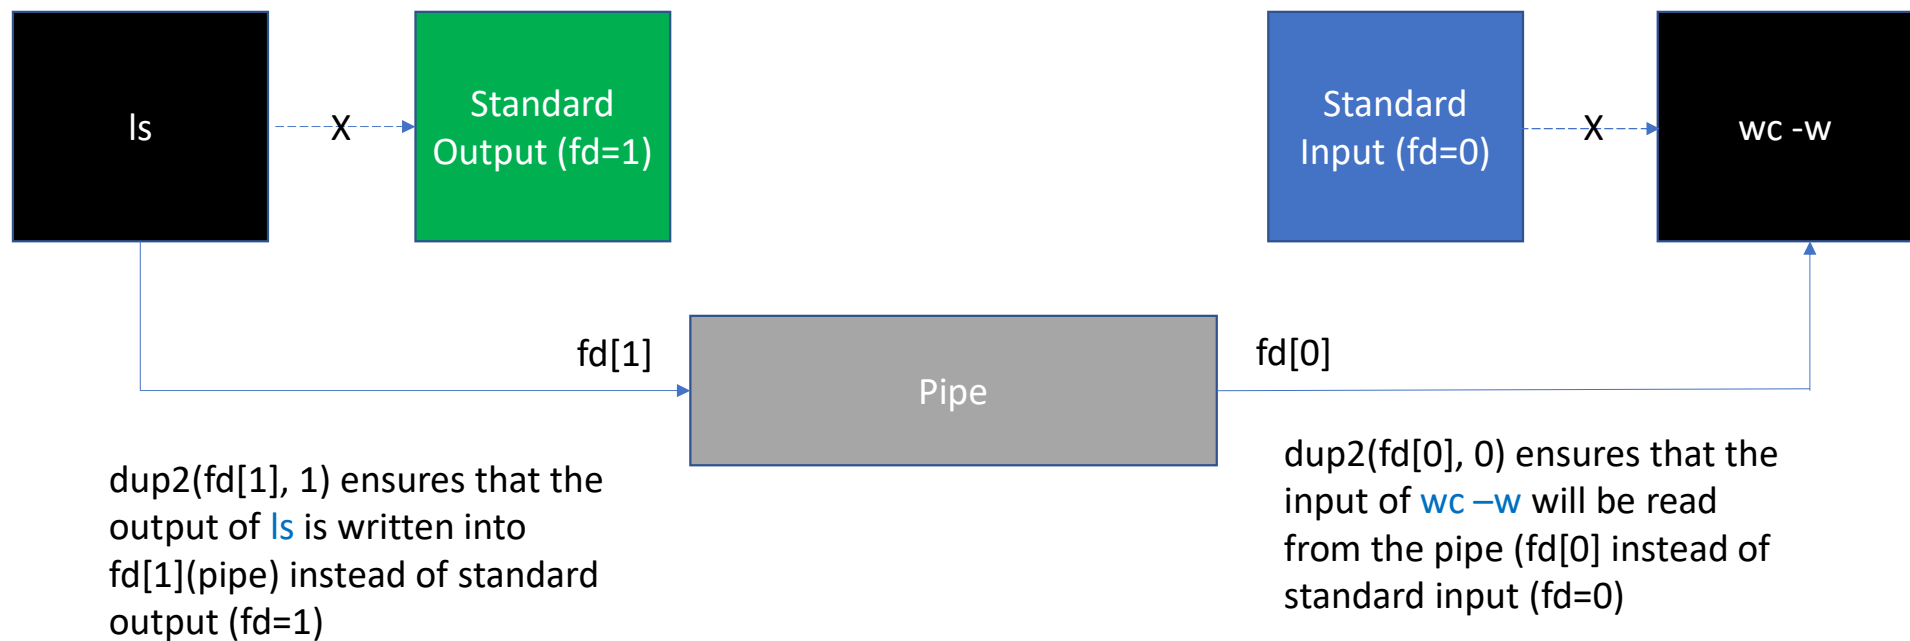
```
int num;
```

```
scanf("%d",&num);
```

```
return EXIT_SUCCESS;
```

```
//end main
```

Implementing the shell pipe mechanism (Example: `$ ls | wc -w`) with `dup2()`



ex4.c //Implementing a shell pipe mechanism

ls | wc -w

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>
//equivalent of $ls|wc -w
int main(int argc, char *argv[]){
    int fd[2];
    if(pipe(fd)==-1)
        exit(1);
    if(fork() > 0) //Parent
    {
        close(fd[0]);
        dup2(fd[1], 1);
        execlp("ls","ls", NULL);
    }
    else //Child
    {
        close(fd[1]);
        dup2(fd[0], 0);
        execlp("wc","wc","-w", NULL);
    }
} //End main
```

FIFOs or named pipes

FIFOs(First In First Out), sometimes called named pipes, offer the following advantages over pipes :

- They have a name that exists in the file system.
- They can **be used by unrelated** processes (//unlike unnamed pipes which can be used by related processes only)
- They exist until explicitly deleted (unnamed pipes exist only as long as the creating process exists)
- FIFO or named pipes are also created in-memory(by the kernel) and are not available in the persistent storage

The system call mkfifo():

```
int mkfifo(const char *path, mode_t mode)
```

mkfifo() returns 0 if OK, -1 otherwise.

Creating a FIFO is similar to creating a file.

Example : `mkfifo("server", 0777);`

- Once a FIFO has been created, it can be treated as a file.
- In particular, the system calls `open()`, `close()`, `read()`, `write()` and `unlink()`(to delete a file) can be used on a FIFO.
- By default, we have :
- Invoking the system call `open()` (for read only) **blocks the caller** until some other process opens the FIFO for writing.
- Invoking the system call `open()` (for write only) **blocks the caller** until some other process opens the FIFO for reading.
- If a process writes to a FIFO that no process has open for reading, **the signal SIGPIPE** will be generated.
- Like pipes, FIFOs are one-way communication channels.
- Note : In case of multiple processes writing to the same pipe, a write of up to `PIPE_BUF` bytes is guaranteed to be atomic (i.e not interleaved)

ex5server.c

```
//A client/server application(within a use and without sockets) where a server (this program)
//accepts data from clients using the FIFO whose name is "server"
//Server program runs on one terminal
//Client program/s run on other terminal/s
#include <fcntl.h>
#include <stdio.h> // This is the server
int main(int argc, char *argv[]){
int fd;
char ch;
unlink("server"); // delete the FIFO file if it exists
if(mkfifo("server", 0777)!=0)//Create the FIFO file
exit(1);
chmod("server", 0777); //there might be a umask
printf("Waiting for a client\n");
fd = open("server", O_RDONLY);
printf("Got a client: ");
// The read call blocks until data is written to the pipe,
// until one end of the pipe is closed,
// or the FIFO is no longer open for writing.
while(read(fd, &ch, 1) == 1)
printf("%c", ch);
} //End main
```

ex5client.c

```
#include <fcntl.h> //Client
#include <stdio.h>

int main(int argc, char *argv[]){
    int fd;
    char ch;
    int count=100;

    while((fd=open("server", O_WRONLY))!=-1)
    {
        printf("trying to connect to the server\n");
        sleep(1);
    }
    printf("Connected: type in data to be sent\n");
    while((ch=getchar()) != -1) // -1 is CTR-D
        write(fd, &ch, 1);

    close(fd);
}
```

Summary

- Unnamed Pipes //Related Processes
- The pipe() system call
- The dup2() system call //To deference file descriptors
 - I/O redirection and reversal
 - Role in implementing the piping command
- FIFOs or named pipes //Can be used by unrelated processes

THANK YOU